



DEPARTMENT OF ELECTRONIC &  
TELECOMMUNICATION ENGINEERING  
UNIVERSITY OF MORATUWA

**Assignment 01 - Intensity  
Transformation and Neighbourhood  
Filtering**

By:

220399B M. M. M. Muftee

Submitted in partial fulfillment of the requirements for the module  
EN3160 Image Processing and Machine Vision

Date: 12<sup>th</sup> of August 2025

**Github Repository:**

[https://github.com/Muftee-Mysan/  
EN3160---Image-Processing-and-Machine-Vision/tree/main/Assignments/  
Assignment\\_01](https://github.com/Muftee-Mysan/EN3160---Image-Processing-and-Machine-Vision/tree/main/Assignments/Assignment_01)

## Contents

1 Question 1: Basic Intensity Transformation	3
2 Question 2: Accentuating Brain Matter	3
3 Question 3: Gamma Correction in L*a*b* Space	3
4 Question 4: Vibrance Enhancement	4
5 Question 5: Histogram Equalization	5
6 Question 6: Foreground Histogram Equalization	6
7 Question 7: Sobel Filtering	7
8 Question 8: Image Zooming	8
9 Question 9: Background Blurring	8

## 1 Question 1: Basic Intensity Transformation

```
# Step 1: Define piecewise transformation function
t1 = np.linspace(0, 50, num=51).astype('uint8')
t2 = np.linspace(100, 255, num=100).astype('uint8')
t3 = np.linspace(150, 255, num=105).astype('uint8')

# Step 2: Concatenate all segments to create the transformation array
t = np.concatenate((t1, t2, t3), axis=0).astype('uint8')
print(t.shape)
```

Figure 1: Code

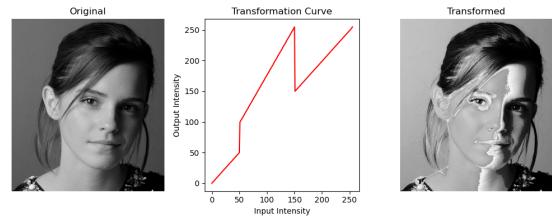


Figure 2: Intensity transformation on emma.jpg image

**Interpretation:** The transformation can be regarded as an identity mapping with a modification that selectively increases the intensity of mid-gray pixels. The key observation is the presence of jump discontinuities in the transformation function, which generate high-contrast regions in the resulting image. In contrast, the high-intensity (white) and low-intensity (dark) pixel values remain unaffected.

## 2 Question 2: Accentuating Brain Matter

**Approach:** A Gaussian function was employed as the intensity transformation function in order to enhance only the desired range of pixel intensities. The choice of a Gaussian function ensures smooth variations in contrast within the output image, thereby avoiding the blocky artifacts that arise from jump discontinuities in the transformation. Furthermore, the mean and standard deviation of the Gaussian can be adjusted, providing precise control over the range of pixel intensities to be emphasized. The transformation function was implemented using the following code:

```
# Step 2: Define the gaussian pulse for white matter
mu = 150
sigma = 20
x = np.linspace(0, 255, 256)
white = 255 * np.exp(-(x - mu)**2) / (2 * sigma**2)

# Step 3: Define the gaussian pulse for gray matter
mu = 200
sigma = 15
x = np.linspace(0, 255, 256)
gray = 255 * np.exp(-(x - mu)**2) / (2 * sigma**2)
```

Figure 3: Code

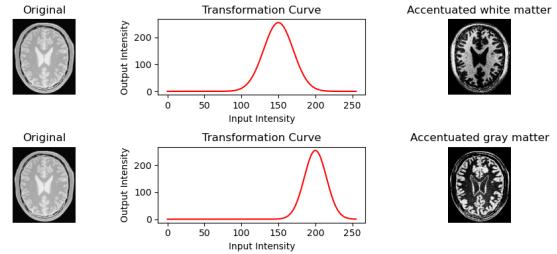


Figure 4: Intensity transformations on brain MRI image

**Interpretation of Results:** The transformation applied to highlight white matter is centered around an intensity value of approximately 150, producing the output image. Similarly, the transformation designed to emphasize gray matter is centered around an intensity value of approximately 200, yielding the output image. It is important to note that the inherent noise present in the original image remains unaffected and is not eliminated by the transformation.

## 3 Question 3: Gamma Correction in L\*a\*b\* Space

**Gamma Correction:** The following code was used to perform gamma correction on the  $L$ -plane of the input image in order to adjust its overall lightness. The OpenCV utility `cv.COLOR_BGR2LAB` was employed to convert the image into the  $L^*a^*b^*$  color space. The gamma value was fine-tuned to  $\gamma = 0.75$ , which effectively lightens the shadow regions while preserving the intensity of the highlights.

$$I_{\text{out}} = I_{\text{in}}^\gamma \quad (1)$$

where  $I_{\text{in}}$  and  $I_{\text{out}}$  denote the input and output pixel intensities, respectively, and  $\gamma$  is the gamma correction parameter.

```
# Split the image in to L*a*b* color space
l,a,b = cv.split(imq3_lab)

# Define and apply gamma transformation
gamma = 0.75
t = np.array([(i/255.0)**gamma*255 for i in np.arange(0, 256)]).astype('uint8')
l_modified = cv.LUT(l, t)

# Merge the channels
merged = cv.merge([l_modified, a, b])
imq3_modified = cv.cvtColor(merged, cv.COLOR_LAB2RGB)
```

Figure 5: Code



Figure 6: Original and gamma corrected output of highlights and shadows.jpg image

**Interpretation of Results:** The histograms of the gamma-corrected image exhibit the presence of noise. Since gamma correction is a non-linear operation, applying it to the  $L$ -plane (which represents lightness in color spaces such as  $L^*a^*b^*$ ) may amplify small variations in pixel intensity. Due to the discrete nature of pixel values, the non-linear transformation redistributes the intensities in a non-uniform manner, thereby introducing quantization noise. As a result, the histogram may appear noisy as pixel intensities are spread unevenly across the range.

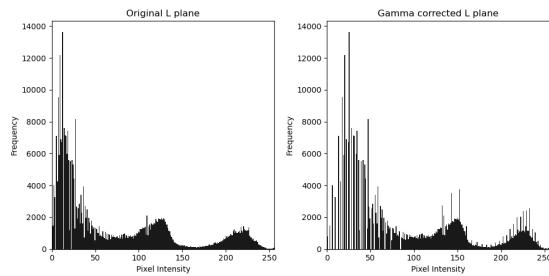


Figure 7: Histograms of the L plane

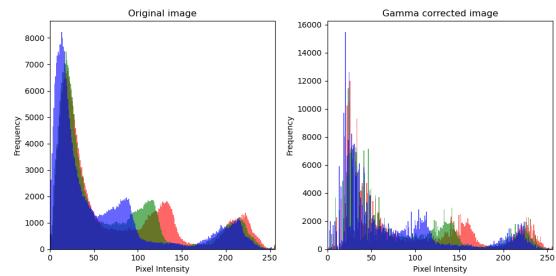


Figure 8: Histograms after converting to RGB

## 4 Question 4: Vibrance Enhancement

### Intensity Transformation in HSV Color Space

Figure 10 illustrates the input image separated into its Hue, Saturation, and Value channels. The following intensity transformation was applied to the Saturation plane:

$$f(x) = \min \left( x + a \times 128 e^{-\frac{(x-128)^2}{2\sigma^2}}, 255 \right) \quad (2)$$

where  $x$  represents the input intensity,  $a \in [0, 1]$ , and  $\sigma = 70$ . By fine-tuning the parameter,  $a = 0.6$  was selected as the optimal value. The resulting transformation is depicted in Figure 6, and the final processed image is shown in Figure 7. The following code snippet was used to apply the transformation exclusively to the Saturation channel.



Figure 9: Image

**Interpretation of Results:** The Saturation plane exhibits high values for the colorful regions of the image (e.g., the red areas of the spider suit), while the colorless or gray regions (e.g., buildings and sky) have low pixel values. By selectively amplifying the intensity of the colorful areas while leaving the colorless regions unchanged, the overall vibrance of the image is enhanced.

```
# Split into planes
H, S, V = cv.split(spider_hsv)
# Define the transformation function
a = 0.6
sigma = 70.0
x = np.arange(0, 256)
F = np.minimum(x + a * 128 * np.exp(-(x - 128)**2) / (2 * sigma**2), 255).astype('uint8')
# Apply transformation to saturation plane
S_modified = cv.LUT(S, f)
# Merge
merged = cv.merge([H, S_modified, V])
spider_modified = cv.cvtColor(merged, cv.COLOR_HSV2RGB)
```

Figure 10: Code

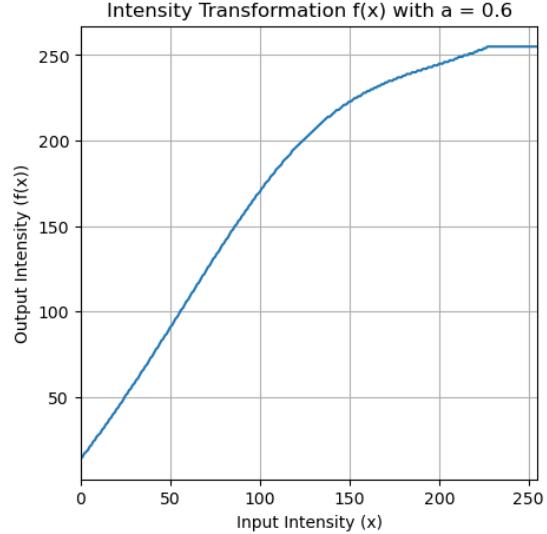


Figure 11: Plot

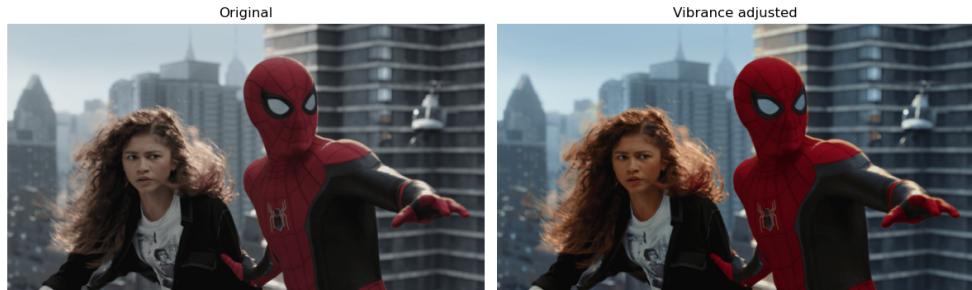


Figure 12: Enhanced Image

## 5 Question 5: Histogram Equalization

```
def histogram_equalization(f):
    # Get image details
    L = 256
    M, N = f.shape

    # Get histogram
    hist = cv.calcHist([f], [0], None, [L], [0, L])
    cdf = hist.cumsum()

    # Define transformation
    t = np.array([(L-1)/(M*N)*cdf[k] for k in range(256)]).astype("uint8")

    return t[f]
```

Figure 13: Code

### Interpretation of Results

As shown in Figure 15, the histogram of the image has become more evenly distributed following histogram equalization. However, this process has also introduced some noise, attributable to quantization effects. Notably, the low-intensity peaks remain largely unchanged.

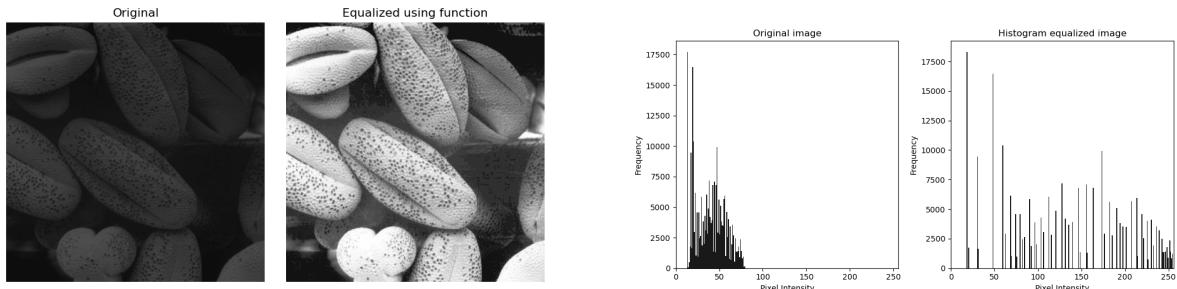


Figure 14: Output image after histogram equalization

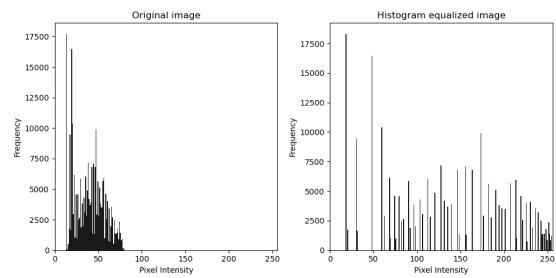


Figure 15: Histograms after equalization

## 6 Question 6: Foreground Histogram Equalization

```
# Apply thresholding on saturation plane
_, mask = cv.threshold(s, 12, 255, cv.THRESH_BINARY)

# Use bitwise_and to extract the foreground using the mask
foreground = cv.bitwise_and(jennifer, jennifer, mask=mask)
```

Figure 16: Code

### Foreground Extraction and Histogram Analysis

Figure 17 illustrates the decomposition of the image `jennifer.jpg` into its Hue, Saturation, and Value channels. The Saturation plane is particularly suitable for generating a mask for foreground extraction. The following code was used to extract the foreground using this mask.

The extracted foreground is shown in Figure 18. The initial histogram of the foreground is presented in Figure 19, while the corresponding cumulative distribution function (CDF) is shown in Figure 20. Figure 21 depicts the histogram after applying histogram equalization to the foreground.



Figure 17: Figure

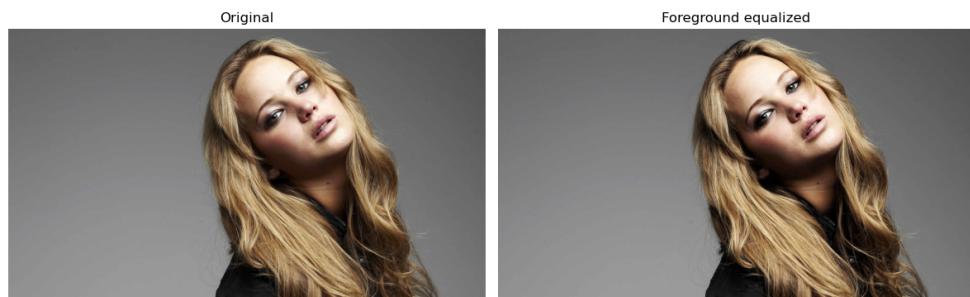


Figure 18: Results of histogram equalization to the foreground of `jennifer.jpg` image

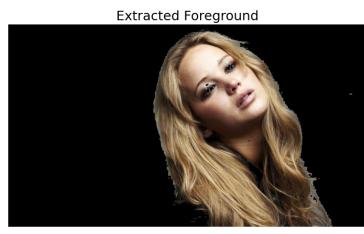


Figure 19: Using cv.filter2D()

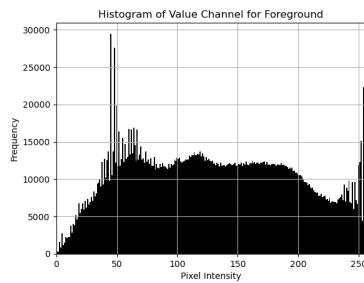


Figure 20: Using custom function

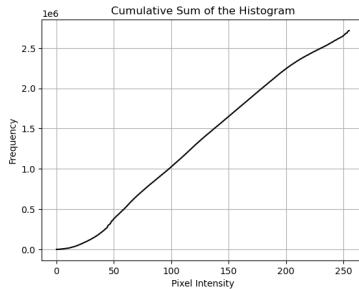


Figure 21: Using cv.filter2D()

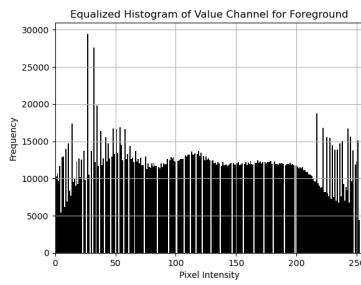


Figure 22: Using custom function

**Interpretation of Results:** The result of histogram equalization applied to the foreground is shown in Figure 22. The equalization process has increased the number of dark pixels, as can be observed by comparing Figures 22(a) and 22(b). This effect is also evident in the output image, where the shadows beneath the face appear darker.

## 7 Question 7: Sobel Filtering

```
def zoom(img, technique, scale=4):
    if technique == 'nn':
        return cv.resize(img, None, fx=scale, fy=scale, interpolation=cv.INTER_NEAREST)
    elif technique == 'bilinear':
        return cv.resize(img, None, fx=scale, fy=scale, interpolation=cv.INTER_LINEAR)
```

Figure 23: Code

### Sobel Edge Detection Using 2D Convolution

The following custom function was implemented to apply a 2D Sobel kernel using 2D convolution. Figure 24 shows the result obtained using OpenCV's `filter2D()` function, while Figure 25 presents the corresponding result produced by the custom function. Figure 26 illustrates the intermediate steps during the separable convolution of the Sobel kernels.

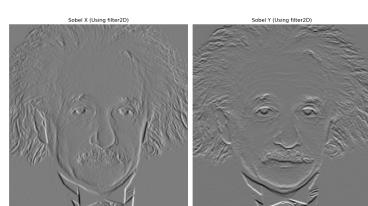


Figure 24: Using cv.filter2D()

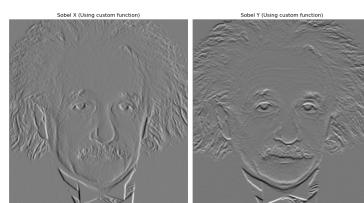


Figure 25: Using custom function

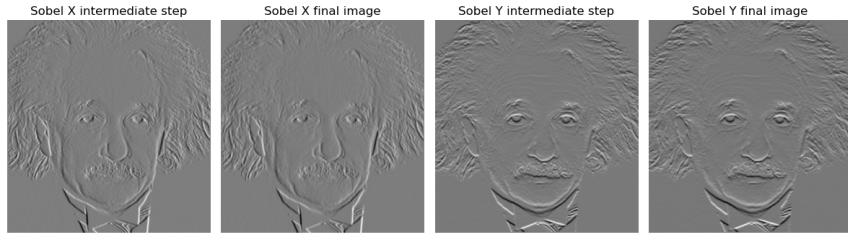


Figure 26: Using Separable Convolution to filter the einstein.png image

## 8 Question 8: Image Zooming

The following function was used to zoom the images using a given interpolation method. The SSD values calculated for each image is given above it.

```
def zoom(img, technique, scale=4):
    if technique == 'nn':
        return cv.resize(img, None, fx=scale, fy=scale, interpolation=cv.INTER_NEAREST)
    elif technique == 'bilinear':
        return cv.resize(img, None, fx=scale, fy=scale, interpolation=cv.INTER_LINEAR)
```

Figure 27: Code



Figure 28: Image 01



Figure 29: Image 02

## 9 Question 9: Background Blurring

```
mask = np.zeros(daisy.shape[:2],np.uint8)
bgdModel = np.zeros((1,65),np.float64)
fgdModel = np.zeros((1,65),np.float64)
rect = (50,100,550,490)
cv.grabCut(daisy,mask,rect,bgdModel,fgdModel,5,cv.GC_INIT_WITH_RECT)
mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
foreground = daisy * mask2[:, :, np.newaxis]
background = cv.subtract(daisy, foreground)
```

Figure 30: The following code utilizes Grabcut to segment the foreground.

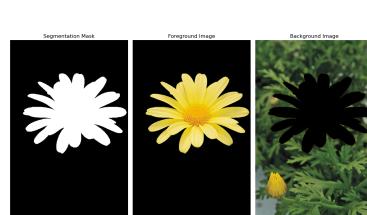


Figure 31: Separation using Grabcut

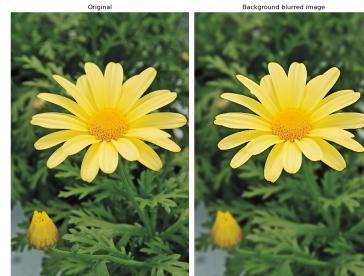


Figure 32: Background blurred

**Reason for the Dark Background Beyond the Edge of the Flower:** When the background image is blurred using a Gaussian kernel, the pixels belonging to the foreground are replaced with zero values. Since Gaussian blur functions as an averaging filter, the regions near the object boundaries are blended with these zero-valued pixels. As a result, the border areas appear darker.